

The Kernel Hacker's Guide to the Galaxy

Automating Exploit
Engineering Workflows

H2HC 2024 São Paulo



whoami (we)



X-FORCE OFFENSIVE RESEARCH



chompie

weird machine mechanic

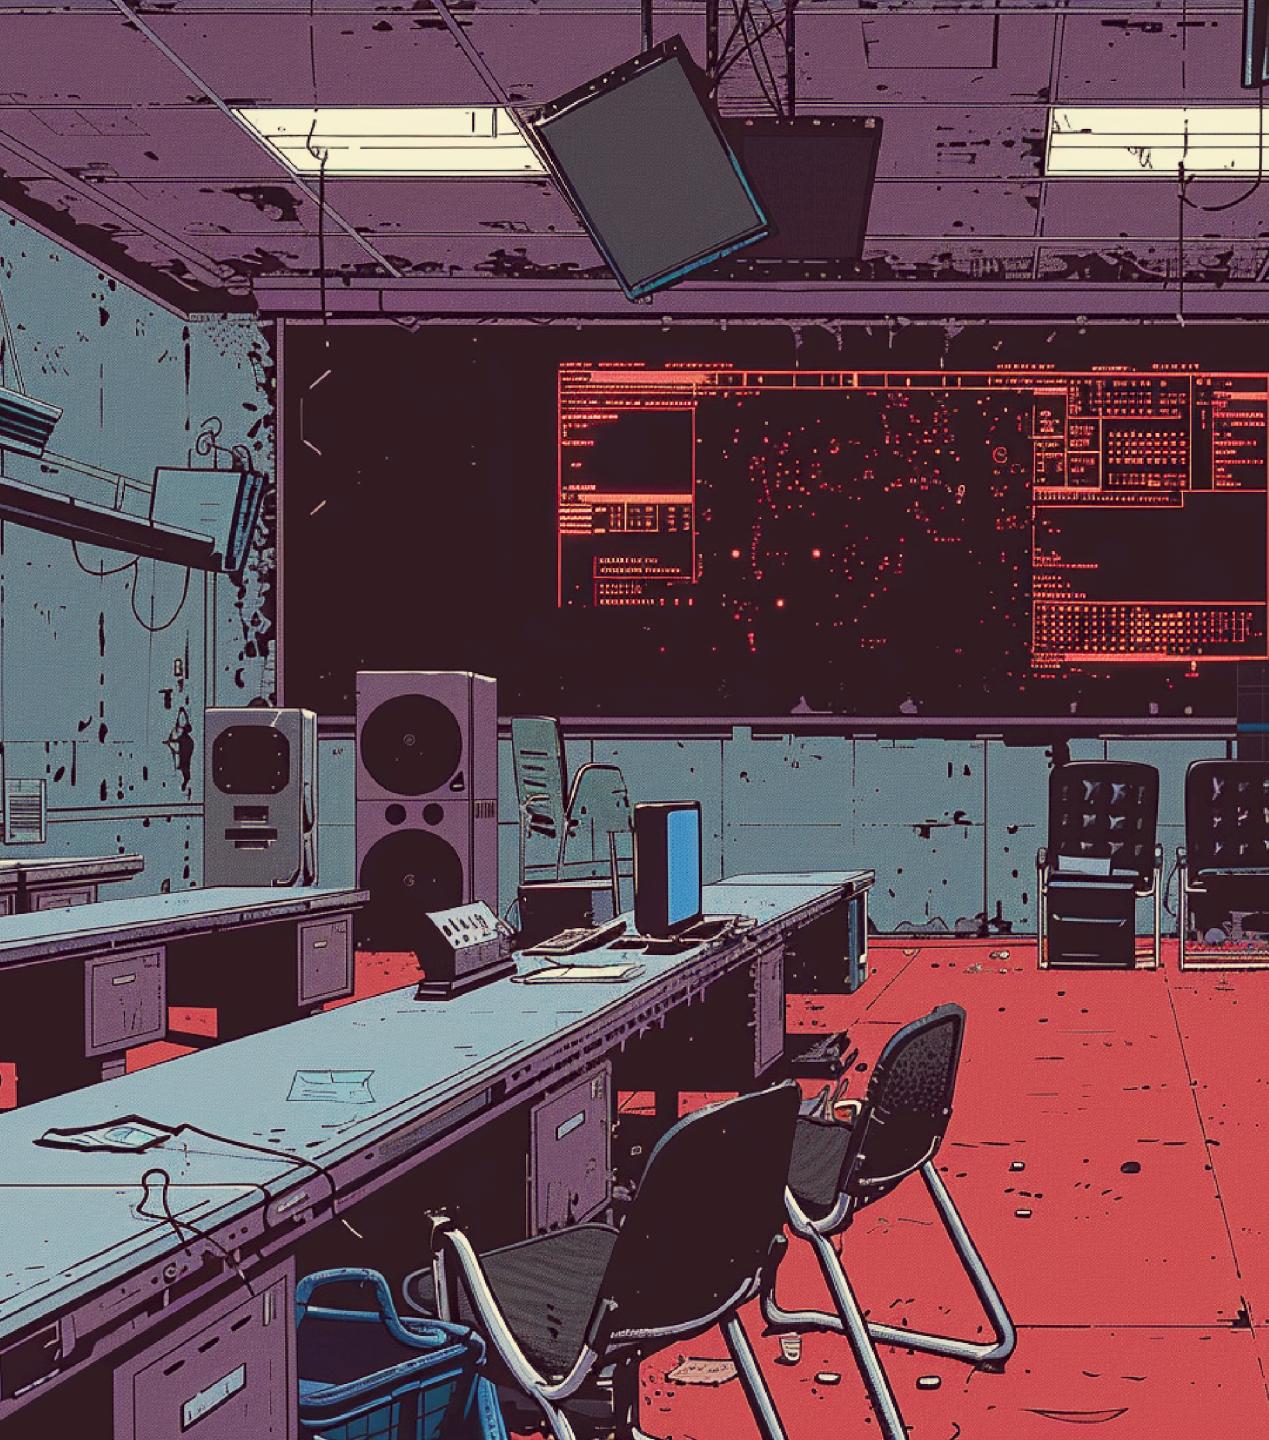
@chompie1337



FuzzySec

just-in-time developer

@FuzzySec



Introduction

Setting the Stage

Why are we talking about this?

- **Practical business need**
 - N-Day **Patch Gaps**
 - **Sustain our sanity**
 - Reduce patch analysis time
 - **Approach a return** to Patch Tuesday ➔ Exploit Wednesday
- **Modern mitigations** and **increased complexity**
 - Exploit construction is becoming **increasingly more labour intensive**
 - **Different targets** have differing considerations & **complexity** (e.g. mobile vs desktop vs IOT)

Exploits are code too

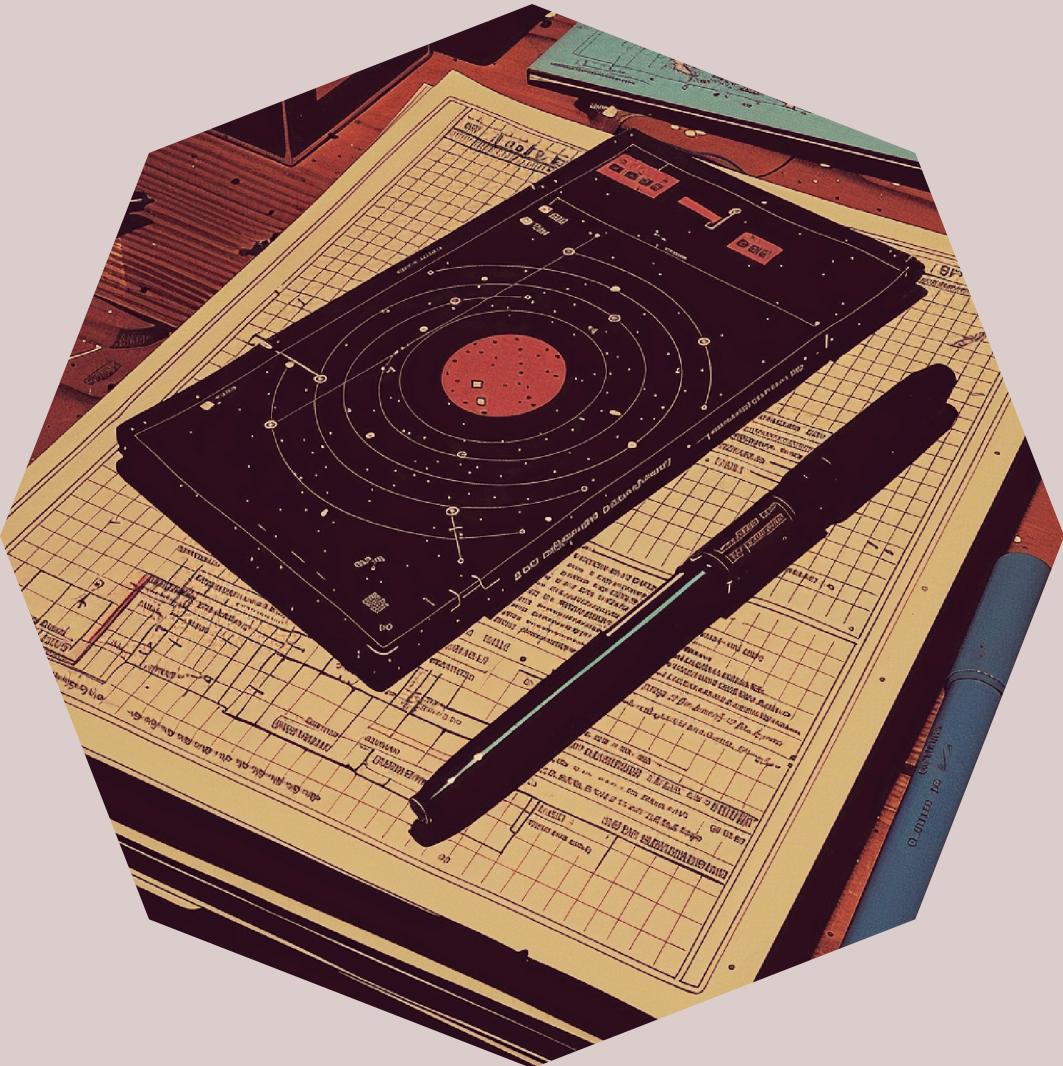
There has been much discussion on **automatic vulnerability discovery**. What about **automating exploit building**, particularly for vulnerabilities that produce **non-trivial primitives**?

Past work

- Practical Uses of Program Analysis: Automatic Exploit Generation
 - Sophia D'Antoine
 - <https://www.youtube.com/watch?v=d3fy4se4JO0>
- The Path Towards Automated Heap Exploitation
 - Thaís Moreira Hamasaki
 - <https://www.youtube.com/watch?v=5XQ0cAwIWMo>
- Greybox Automatic Exploit Generation for Heap Overflows in Language Interpreters
 - Sean Heelan
 - <https://www.cs.ox.ac.uk/tom.melham/phd/Heelan-2020-GAE.pdf>
- Hot topic in 2016 Darpa Cyber Grand Challenge
 - Focus shifted now towards AI vulnerability discovery



Theory & Practice



- Past work is **largely theoretical**
- Focus on applying Automated Exploit Engineering (**AEG**) concepts to real targets
 - Windows Kernel LPE exploits
- **Historical** thought process
 - AEG will mostly apply to **granular workflows**
 - This **aligns with our efforts**

Theoretical AEG

Two Phases

- Code to **trigger** the vulnerability
- Generate **payload** from weird machine **primitive**

AEG in Practice

Many Phases

- Code to **trigger** the vulnerability
- Exploit **primitive construction/transformation**
- **Mitigation bypasses**
 - **Information leaks, CFI gadgets, additive primitive contortion**, avoidance of Kernel Patch Protection (**KPP**) and **HVCI**
- **Stability** Improvements
- Targeting
- Stealth

The Exploit Development Lifecycle (BSides Canberra)

Keynote 2024) by **chompie**

- <https://drive.google.com/file/d/1jHnVdjAcPGkuVPiakZBAOTp8uzMej6LY/view>

What is automation, really?



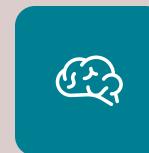
Immediate practical impact

- Self contained
- **Atomic operations**



Reduce manual labour

- Primitive discovery is an intensive task
- Suitable **exploitation objects**
- Control Flow **gadgets**



Maintain sanity

- **Narrow search-time** for broad deterministic questions
- **Save brainpower** for reverse engineering tasks that can't be automated (yet)



Maintenance & QA

- **Does this work** on this other targets?

Component Breakdown

1

Initial Trigger

- Generate **user-space code** to trigger the vulnerability

2

Exploitability analysis

- Code **reachability**
- Race **condition analysis**

3

Exploit primitive construction

- Heap spray **object identification**

4

Mitigation bypasses

- **Primitive contortion**
- kCFG gadgets

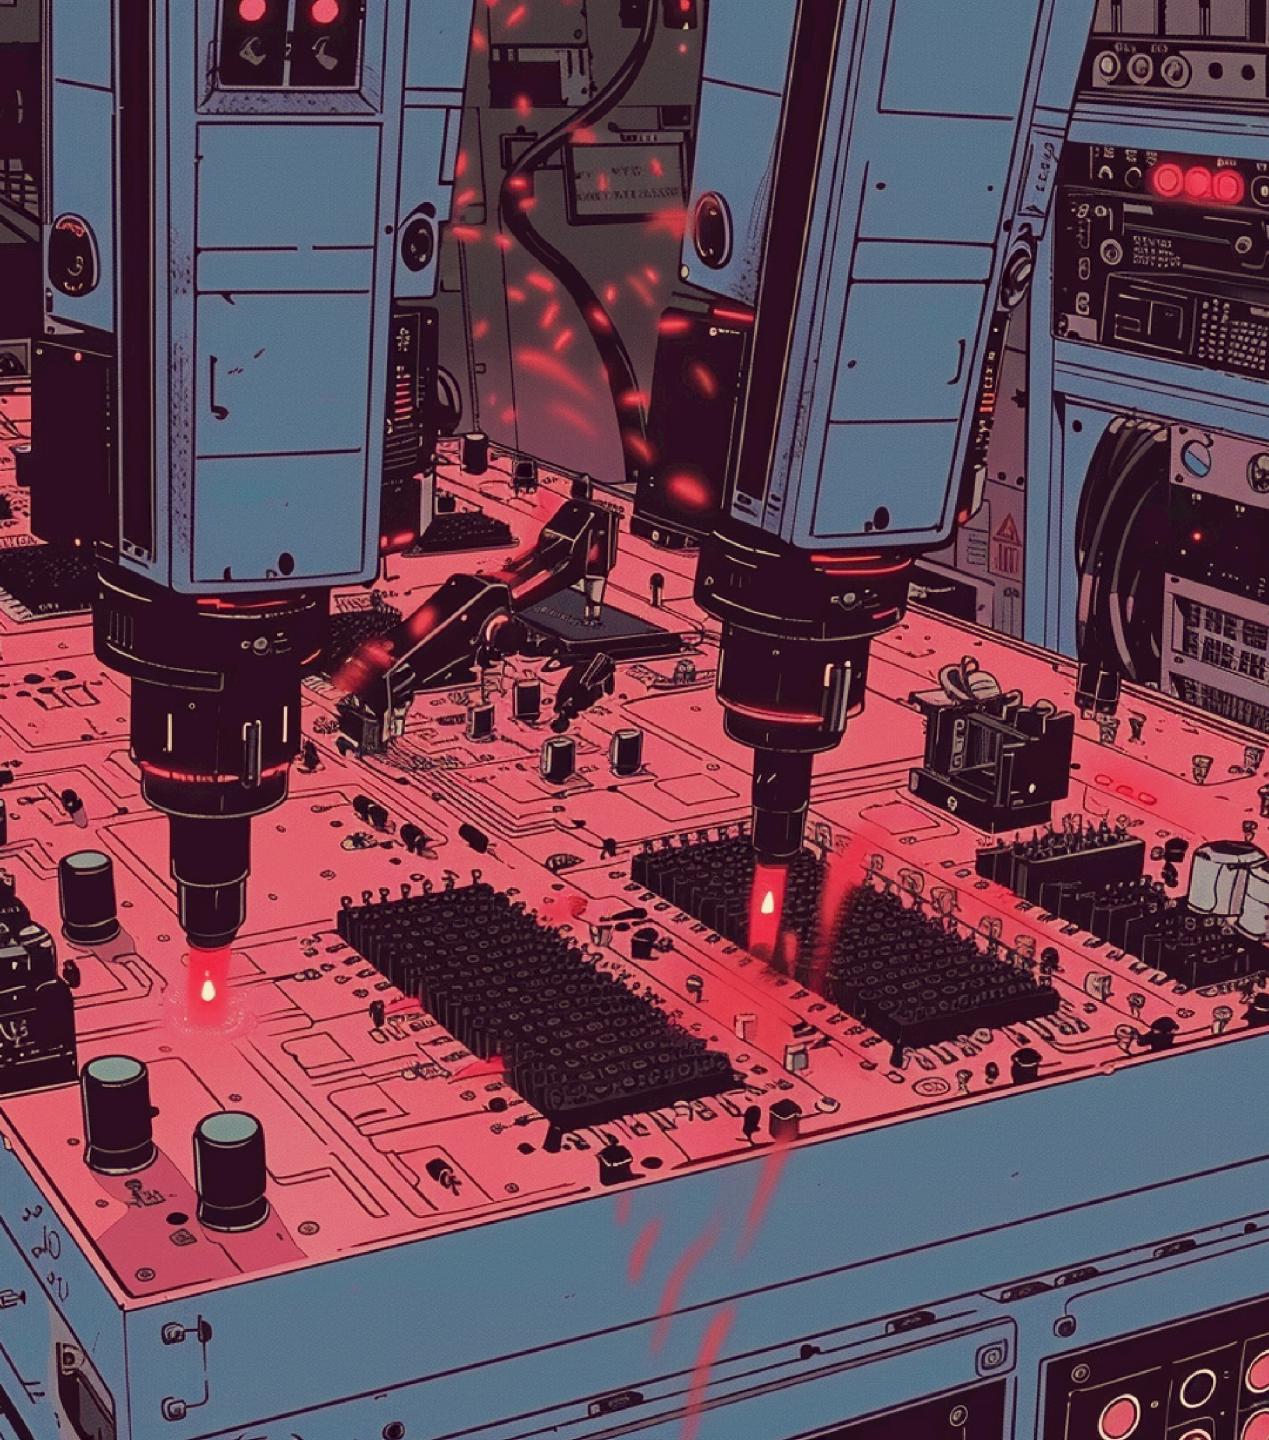
5

Stability Improvements

6

Variant analysis

- Does this pattern exist somewhere else?

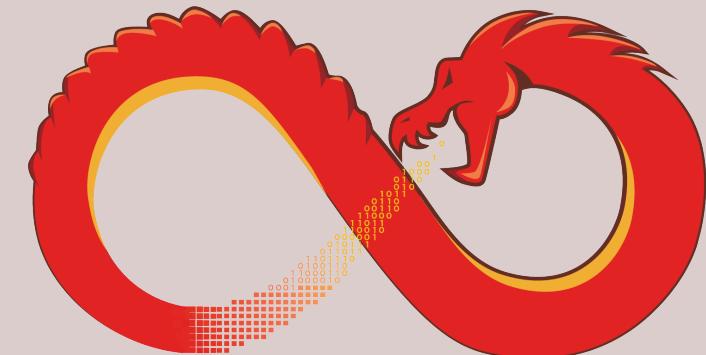


The Atomics of Automation

Program Analysis

Choose your weapon(s)

- Research requires a toolchain
 - Base tool capability is important (architectures, decompilers, quality)
 - Vulnerability research requires building tools on top of your toolchain, tool choices cascade into your entire pipeline
 - Support for automation, headless processing and API is critical
- XOR is primarily a Binary Ninja shop now, with redundancy using Ghidra
 - Vector35 has been a good partner for reporting bugs and extending Ninja capabilities
 - Decompiler is good and has been appreciably improving over time
 - API support is excellent
 - <https://api.binary.ninja/>
 - More API bindings needed (nodeJS & .NET)



GHIDRA

Duck Tape

- What is Duck Tape?

- Duck Tape is a **binary diffing engine**
- **Monaco** editor **code visualization**
- **Currently** due to patch ingestion it is **focused on Windows** but the diffing **engine generalizes** to other platforms

- Why do you need a **diffing engine**?

- We want to **reduce cognitive-load**
- Operating system **patches are vast** and it is not always possible to **track and inspect all changes**
- **Patch logs don't provide granular information** and also omit changes that have been pushed
- A lot of inspiration from **@TinySecEx** ❤️

Status

Worker status for Duck Tape backend and queue manager.

NAME	IP	PORT	STATUS
binja-1	10.1.1.10	8000	Online
binja-2	10.1.1.11	8000	Online

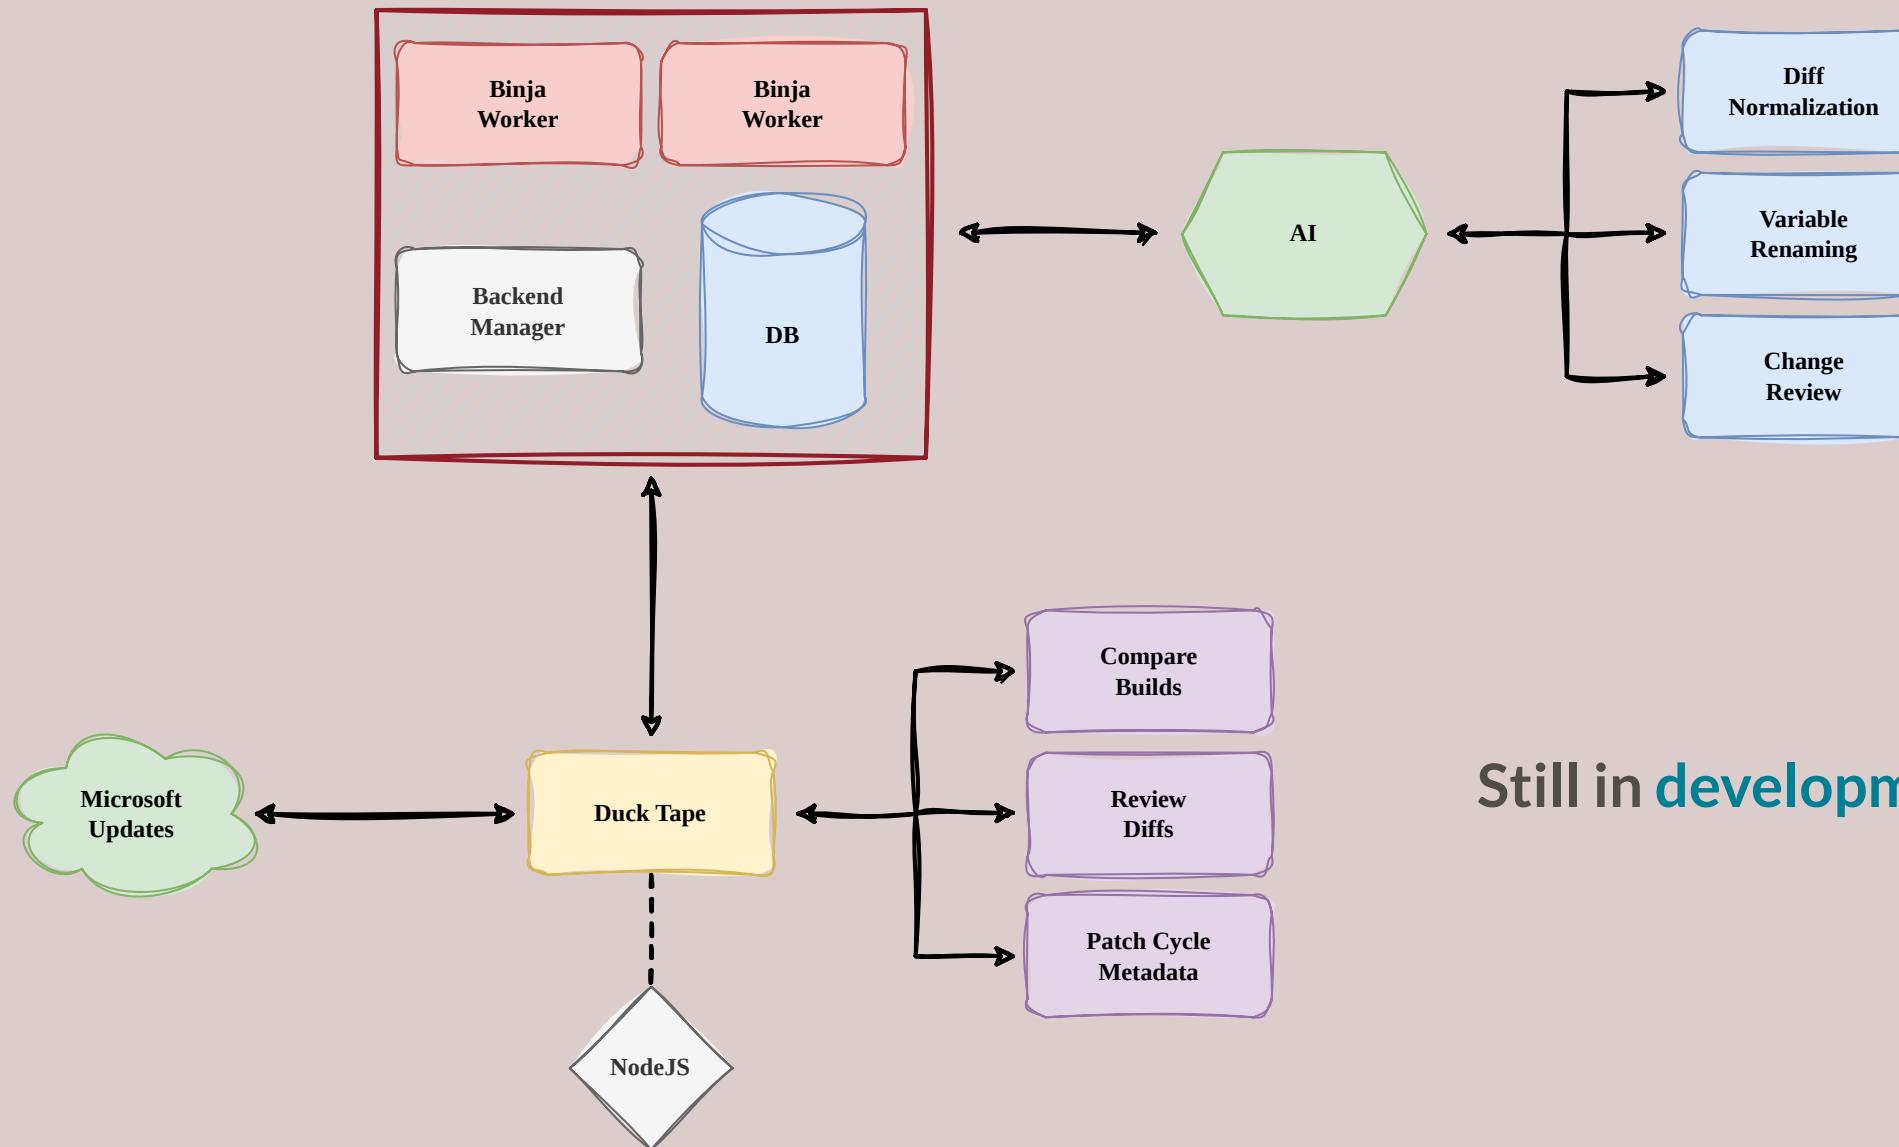
Target Platform Date

11-23H2

Queue Job

Close

Theory Of Design



Still in **development!**



Diff for tcpip.sys (10.0.22621.3958 → 10.0.22621.4036)

tcpip

FILEN

net1

net1

tcpip

tcpip

tcpip

tcpip

```
183 |     *(arg1 + 0x2c) = 0x3b
184 |     int32_t var_60_1 = 0
185 |     var_68.q = 0
186 |     *(arg1[1] + 0x8c) = 0xc000021b
187 |
188 |
189- if (IppDiscardReceivedPackets(&Ipv6Global, result_1, arg1,
190     int32_t var_58_1
191     var_58_1.b = r14.b
192 |
193- IppSendErrorList(1, &Ipv6Global, arg1, 4, var_68.b, _bs
194 |
195     result = zx.q(result_1)
196     label_1c006dc78:
197     __security_check_cookie(rax_1 ^ &var_88)
198     return result
```

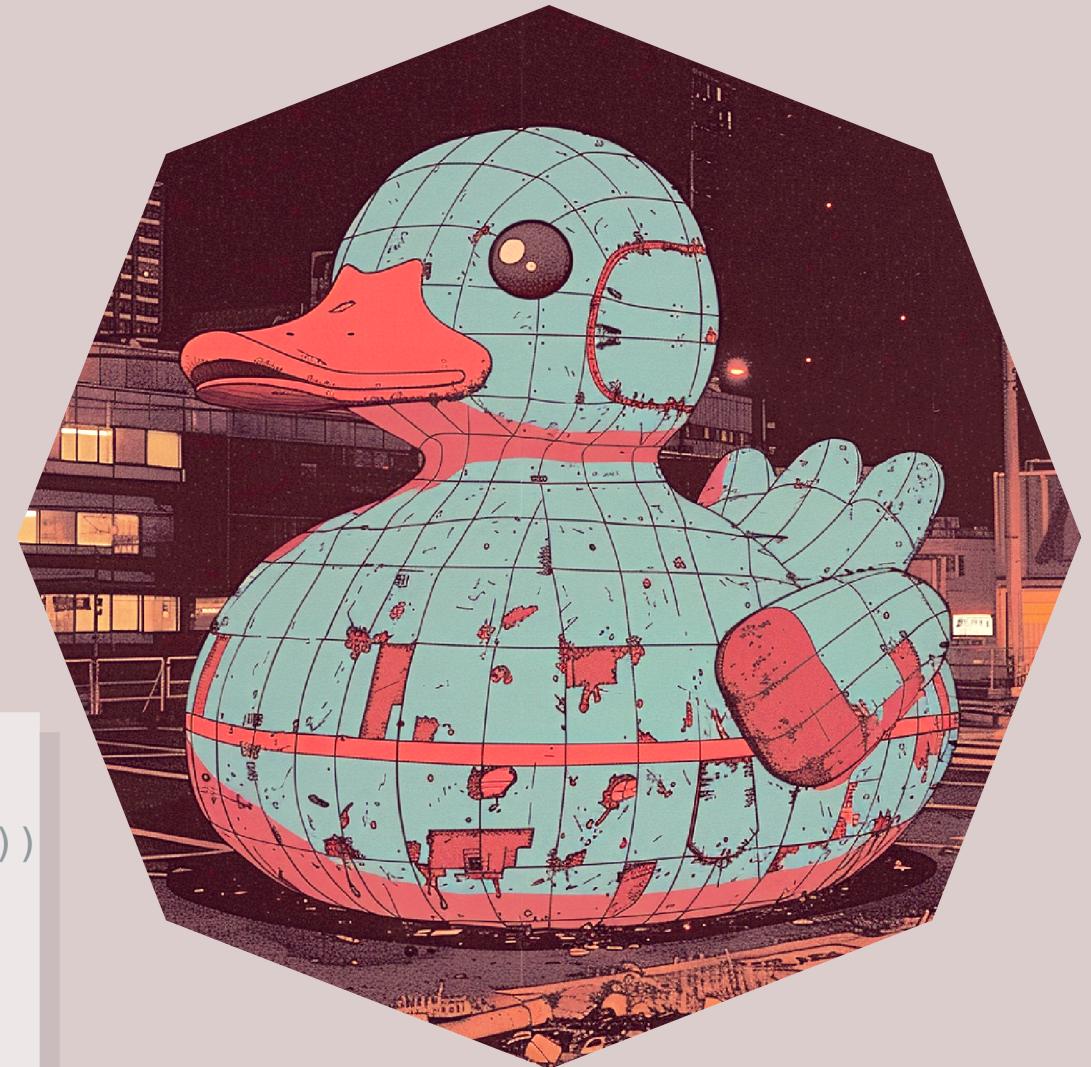
```
183 |     *(arg1 + 0x2c) = 0x3b
184 |     int32_t var_60_1 = 0
185 |     var_68.q = 0
186 |     *(arg1[1] + 0x8c) = 0xc000021b
187 |
188 |
189+ if (IppDiscardReceivedPackets(&Ipv6Global, result_1, arg1,
190     int32_t var_58_1
191     var_58_1.b = r14.b
192+
193+
194+
195+
196+
197     var_68.b = r12.b
198+
199+
200+
201+
202+
203
204
205
206
207
```

[Close](#)

Automating Root Cause Analysis (RCA)

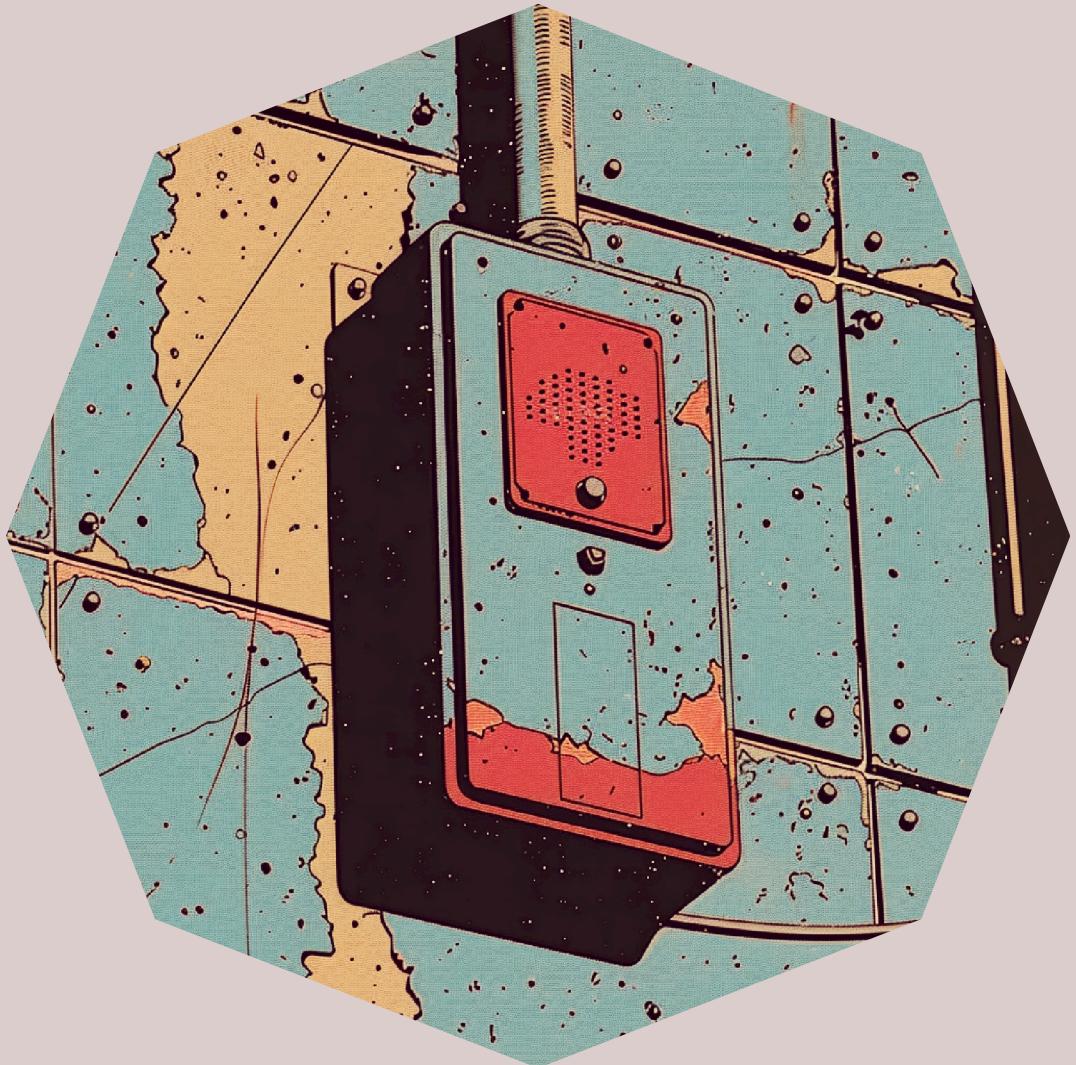
- Programmatically **pull** and **diff** patches
- **Tell-tale markers** for security issues
 - All **Microsoft patches** are introduced via **feature flag**
 - Easy to identify
 - Clean **condition** to use for **RCA**
 - If Feature Flag enabled, patched code path
 - Else vulnerable code path
- **CVE-2024-38063** tcpip RCE ([@XiaoWei__](#) at [Kunlun Lab](#))

```
if  
(!Feature_1223761209_private_IsEnabledDeviceUsage())  
    IppSendErrorList(rcx_17, &Ipv6Global, arg1,  
r9_2, var_68, temp0_3, var_58_1);  
else  
    IppSendError(rcx_17, &Ipv6Global, arg1, r9_2,  
var_68, temp0_3, var_58_1);
```



Trigger Mechanisms

- Vulnerable code paths chain up from user mode reachable interfaces using program analysis toolkit
 - NT Kernel & Win32k ⇒ system calls
 - Drivers ⇒ DispatchDeviceControl, DispatchRead, DispatchWrite, and DispatchReadWrite, DispatchCleanup, DispatchCreate, etc.
- Select vulnerable function component and **generate condensed backtrace** based on cross-references
 - Search upwards till **user entrypoint** is found
 - Feed condensed path **to LLM worker** for POC generation (maybe)



Case-Study CVE-2024-21338

- From Lazarus (North Korea) with love
- There are actually **two bugs** here that can be exploited, **not one**
 - **AipAllocateImageContext** has an **arbitrary increment** primitive
 - **AppHashComputeImageHashInternal** has an **arbitrary call** primitive

```
Breakpoint 0 hit
appid!AppHashComputeImageHashInternal+0x78:
fffff804`5b53e18c 488b00      mov    rax,qword ptr [rax]
2: kd> r
rax=4343434343434343 rbx=0000000000000000 rcx=fffffe50e8821eb00
rdx=0000000000000000 rsi=0000000000000000 rdi=fffff958de06d15ec
rip=fffff8045b53e18c rsp=fffff958de06d1370 rbp=fffff958de06d1409
r8=0000000000000001 r9=fffff958de06d14c0 r10=fffff80455166a30
r11=fffff958de06d13a0 r12=0000000000000001 r13=0000000000000000
r14=000000000000000c r15=0000000000000000
iopl=0    nv up ei pl nz na pe nc
cs=0010 ss=0018 ds=002b es=002b fs=0053 gs=002b          efl=00040202
appid!AppHashComputeImageHashInternal+0x78:
fffff804`5b53e18c 488b00      mov    rax,qword ptr [rax] ds:002b:43434343`43434343=?????????????????
2: kd> u
appid!AppHashComputeImageHashInternal+0x78:
fffff804`5b53e18c 488b00      mov    rax,qword ptr [rax]
fffff804`5b53e18f 488d55ff    lea    rdx,[rbp-1]
fffff804`5b53e193 ff1577b5feff  call   qword ptr [appid! guard dispatch icall fptr (fffff804`5b529710)]
```

```
2: kd> r
rax=fffffa802f2884ed0 rbx=0000000000000001 rcx=4242424242424242
rdx=fffffa802f2884f98 rsi=4242424242424242 rdi=0000000000000000
rip=fffff80738e7e044 rsp=fffffce8e4425a4e0 rbp=0000000000000000
r8=0000000000000000 r9=fffffa802f2884f70 r10=fffff80738e7e020
r11=fffffa802f2d27fb8 r12=0000000000000002 r13=0000000000000000
r14=fffffce8e4425a5a0 r15=fffffe08c33569180
iopl=0    nv up ei pl nz na po nc
cs=0010 ss=0018 ds=002b es=002b fs=0053 gs=002b          efl=00040246
nt!ObfReferenceObject+0x24:
fffff807`38e7e044 f0480fc15ed0    lock xadd qword ptr [rsi-30h],rbx ds:002b:42424242`42424212
2: kd> k
# Child-SP           RetAddr          Call Site
00  fffffce8e`4425a4e0  fffff807`419c97dd  nt!ObfReferenceObject+0x24
01  fffffce8e`4425a520  fffff807`419ca4e4  appid!AipAllocateImageContext+0xb9
```



Automate all the things!

But how do we reach these code paths?

```
PS C:\Users\b33f\tools\binja-scripts\symbol-searcher> python.exe .\backtrace.py c:\Users\b33f\tools\ExpDev\CVE-2024-21338-APPID\appid.sys-10.0.22621.2506.b  
lob -s AppHashComputeImageHashInternal  
[+] Starting backtrace for function: AppHashComputeImageHashInternal (0x1C002E114)  
  
Total input handlers found: 0  
Total Call Paths Identified: 4  
Total Unique Functions in Call Paths: 19  
  
[+] Call Trace (Shortest Paths First):  
  
AppHashComputeImageHashInternal -> AppHashComputeFileHashesInternal -> AipSmartHashImageFile -> AipDeviceIoControlDispatch -> DriverEntry  
[Input Handlers:  
Address: 0x1C002E114 | Prototype: uint64_t AppHashComputeImageHashInternal(int64_t arg1, int64_t* arg2, int32_t arg3, int64_t* arg4)  
Address: 0x1C0002C94 | Prototype: uint64_t AppHashComputeFileHashesInternal(int64_t arg1, int64_t* arg2, int64_t arg3, uint128_t* arg4, int64_t arg5)  
Address: 0x1C002A46C | Prototype: uint64_t AipSmartHashImageFile(void* arg1, HANDLE arg2, void* arg3, void* arg4)  
Address: 0x1C001DF20 | Prototype: uint64_t AipDeviceIoControlDispatch(int64_t arg1, IRP* arg2)  
Address: 0x1C00350C4 | Prototype: uint64_t DriverEntry(DRIVER_OBJECT* arg1)  
]
```

```
PS C:\Users\b33f\tools\binja-scripts\symbol-searcher> python.exe .\backtrace.py c:\Users\b33f\tools\ExpDev\CVE-2024-21338-APPID\appid.sys-10.0.22621.2506.b  
lob -s AipAllocateImageContext  
[+] Starting backtrace for function: AipAllocateImageContext (0x1C0029724)  
  
Total input handlers found: 0  
Total Call Paths Identified: 4  
Total Unique Functions in Call Paths: 15  
  
[+] Call Trace (Shortest Paths First):  
  
AipAllocateImageContext -> AipSmartHashImageFile -> AipDeviceIoControlDispatch -> DriverEntry  
[Input Handlers:  
Address: 0x1C0029724 | Prototype: uint64_t AipAllocateImageContext(int64_t arg1, HANDLE arg2, int64_t arg3, void*** arg4)  
Address: 0x1C002A46C | Prototype: uint64_t AipSmartHashImageFile(void* arg1, HANDLE arg2, void* arg3, void* arg4)  
Address: 0x1C001DF20 | Prototype: uint64_t AipDeviceIoControlDispatch(int64_t arg1, IRP* arg2)  
Address: 0x1C00350C4 | Prototype: uint64_t DriverEntry(DRIVER_OBJECT* arg1)  
]
```

What about CVE-2024-30088?

- Exploited by Emma Kirkpatrick (@carrot_c4k3) at Pwn2Own 2024 in Vancouver
 - Used ITW by APT34 (Iran)
 - TOCTOU leading to arbitrary write using `RtlCopyUnicodeString`
 - Reachable through `NtQueryInformationToken` \Rightarrow `TokenAccessInformation`

```
[+] Starting backtrace for function: AuthzBasepCopyoutInternalSecurityAttributes (0x140687EA0)

Total input handlers found: 0
Total Call Paths Identified: 4
Total Unique Functions in Call Paths: 6

[+] Call Trace (Shortest Paths First):

AuthzBasepCopyoutInternalSecurityAttributes -> AuthzBasepQueryInternalSecurityAttributesToken -> SepGetTokenAccessInformationBufferSize -> NtQueryInformationToken
    [Input Handlers:
        Address: 0x140687EA0 | Prototype: uint64_t AuthzBasepCopyoutInternalSecurityAttributes(int32_t* arg1, uint128_t* arg2, int32_t arg3)
        Address: 0x140687D74 | Prototype: int64_t AuthzBasepQueryInternalSecurityAttributesToken(int32_t* arg1, uint128_t* arg2, int32_t arg3, int32_t* arg4)
        Address: 0x140687B78 | Prototype: uint64_t SepGetTokenAccessInformationBufferSize(void* arg1, char arg2, void* arg3, int32_t* arg4, int32_t* arg5, int32_t* arg6, int32_t* arg7, int32_t* arg8, int32_t* arg9, int32_t* arg10, int32_t* arg11, int32_t* arg12, int32_t* arg13)
        Address: 0x1406EE5A0 | Prototype: uint64_t SeQueryInformationToken(uint128_t* arg1, uint32_t arg2, int32_t* arg3)
    ]
AuthzBasepCopyoutInternalSecurityAttributes -> AuthzBasepQueryInternalSecurityAttributesToken -> SepCopyTokenAccessInformation -> NtQueryInformationToken
    [Input Handlers:
        Address: 0x140687EA0 | Prototype: uint64_t AuthzBasepCopyoutInternalSecurityAttributes(int32_t* arg1, uint128_t* arg2, int32_t arg3)
        Address: 0x140687D74 | Prototype: int64_t AuthzBasepQueryInternalSecurityAttributesToken(int32_t* arg1, uint128_t* arg2, int32_t arg3, int32_t* arg4)
        Address: 0x140687B8C | Prototype: uint64_t SepCopyTokenAccessInformation(void* arg1, int32_t** arg2, int32_t arg3, int32_t arg4, int32_t arg5, int32_t arg6, int32_t arg7, int32_t arg8, int32_t arg9, int32_t arg10, int32_t arg11, int32_t arg12, int32_t arg13, char arg14, int64_t* arg15)
        Address: 0x14069B730 | Prototype: int64_t NtQueryInformationToken(uint64_t arg1, uint32_t arg2, char* arg3, int32_t arg4, int32_t* arg5)
    ]
```

Execution links
up to user
callable entry
point

LLM automation

- Smarter models can do smarter workloads
 - The OpenAI o1 model can actually analyse this backtrace by ingesting function bodies for intermediate calls
 - LLM's with vector stores containing specialized knowledge and RAG can do an ok job at analysis
 - Models can break down tasks into component parts and use function calling in the decompiler to answer their own questions
- LLM's in VR is a critical area that many people are looking at (quietly)
 - Project Naptime: Evaluating Offensive Security Capabilities of Large Language Models (Google P0)
 - From Naptime to Big Sleep: Using Large Language Models To Catch Vulnerabilities In Real-World Code (Google P0)
 - Automated LLM Bugfinders (Dave Aitel)

High-level Control Flow (Only Pertinent Steps)

- NtQueryInformationToken(...) is invoked with TokenInformationClass = TokenAccessInformation (0x16).
- Inside NtQueryInformationToken, it references the token and gathers buffer sizes.
- SepCopyTokenAccessInformation is eventually called to build the combined structures.
- SepCopyTokenAccessInformation calls AuthzBaseQueryInternalSecurityAttributesToken to handle internal security attributes.
- AuthzBaseQueryInternalSecurityAttributesToken calls AuthzBaseCopyoutInternalSecurityAttributes to output the attributes.
- AuthzBaseCopyoutInternalSecurityAttributes invokes RtlCopyUnicodeString to copy the attribute's Unicode strings.

Minimal C Code Example (Skeleton)



```
#include <windows.h>
#include <winternl.h>

// Assume we already have a valid token handle with TOKEN_QUERY access.
// For demonstration, hToken is an already opened token handle.
HANDLE hToken;
DWORD returnLength;
BYTE buffer[1024]; // Sufficiently large buffer

NTSTATUS status = NtQueryInformationToken(
    hToken,
    TokenAccessInformation,      // Trigger the path
    buffer,
    sizeof(buffer),
    &returnLength
);
```



Spraying the Windows Kernel Pool



- Heap/Pool spraying **exploitation technique** for exploiting several common **vulnerability classes**
 - Use after free, Double Free, Out-of-Bounds
- Spray objects must match **vulnerable object** on
 - Paged or NonPaged
 - Size or relative (LFH, Large Pool Alloc, etc)
- **Not much research** on **new objects** for pool spraying on Windows
 - Current methods largely meet requirements in most cases
 - **NonPaged** Pool: Named Pipes (buffered/unbuffered)
 - **Paged** Pool: PipeAttributes, WNF_STATE_DATA
 - Elastic **data size** and user-controlled **data buffer**
- **Sometimes** vulnerability primitives require more work
 - Known spray methods do not apply
 - Influenced/exacerbated by **Kernel mitigations & EDR**
 - **Specialized spray gadget** required

Pool Spray Search (Constant Size)



- Methodology

- Search for calls to allocation routines
- Check **pool flags** match search constraint
- Check size parameter is **HLIL_CONST**
 - Check if **size** parameter is **within range**
- Trace call up to match to a user mode **entry point** to ensure allocation can be triggered from **user space**

- Example gadgets

- Searching **ntoskrnl.exe**
- Searching **POOL_PAGED**
- Varying max size **32, 160, 240**

```
[+] Calling Function --> NtSetUuidSeed
|_ HLIL representation: ExAllocatePool2(0x100, 0x20, 0x64695555)
|_ Address: 0x140846610
|_ Pool Tag: UUId
|_ Size: 0x20
|_ Pool Flags: POOL_FLAG_PAGED

[+] Calling Function --> NtSetInformationJobObject
|_ HLIL representation: ExAllocatePool2(0x100, 0x88, 0x624a7350)
|_ Address: 0x1407d5c75
|_ Pool Tag: PsJb
|_ Size: 0x88
|_ Pool Flags: POOL_FLAG_PAGED

[+] Calling Function --> NtMapCMFModule
|_ HLIL representation: ExAllocatePool2(0x100, 0xf0, 0x636d6650)
|_ Address: 0x140a01994
|_ Pool Tag: Pfmc
|_ Size: 0xf0
|_ Pool Flags: POOL_FLAG_PAGED
```

Pool Spray Search (Elastic Size)



Named Pipes (Buffered)

```
[+] Calling Function --> NpAddDataQueueEntry
|_ HLIL representation: ExAllocatePool2(zx.q((rbx_1 + 0x41).d), zx.q(rcx), 0x7246704e)
|_ Address: 0x1c000c487
|_ Pool Tag: NpFr
|_ Size: HLIL_ZX
[+] Memory copy HLIL representation: memcpy(&rbx_1[6], UserBuffer, zx.q(arg6))
```

Known primitive

- Easy to find other **non-public spray primitives**, elastic in size and user controllable data. Some examples:
 - **NtSetInformationFile**
 - **NtSetVolumeInformationFile**
 - **NtSet EaFile**

- Methodology

- Same as constant search, **except size** is **not a constant**
- Process the remainder of the function
 - **Memory copy** within destination returned by **allocation**?
- Possible to **tighten constraints** (e.g. copy should happen at start of allocation)

Unknown primitive

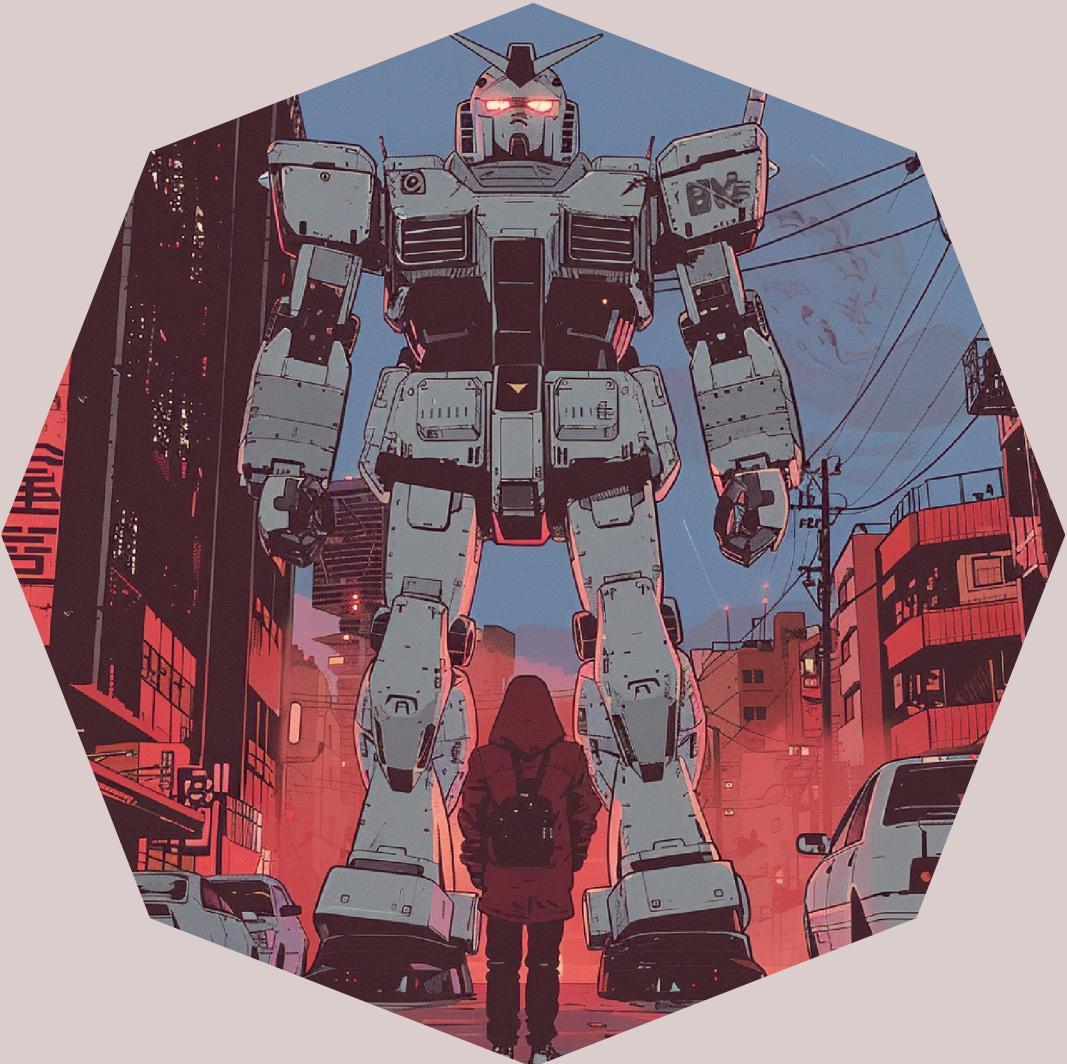
```
[+] Calling Function --> NtSetInformationFile
|_ HLIL representation: ExAllocatePool2(0x61, arg4, 0x42536f49)
|_ Address: 0x1402f0373
|_ Pool Tag: IoSB
|_ Size: HLIL_VAR
|_ Pool Flags: POOL_FLAG_USE_QUOTA, POOL_FLAG_RAISE_ON_FAILURE, POOL_FLAG_NON_PAGED
[+] Memory copy HLIL representation: memcpy(rax_31, arg3, arg4)
```

k(ernel)CFG Primer

- kCFG is an **exploitation mitigation** mechanism in the Windows Kernel that **restricts where execution flow can go**
- Provides protection against techniques like **ROP** - restricts the targets of **indirect calls**
- Like CFG in Windows user mode
- System process has its own **bitmap**, protected by Hypervisor-Protected Code Integrity (**HVCI**)



Mitigations managed



- Mitigation bypass constraints are often well understood. **What is the problem?**
 - **Search space** is **prohibitively large**, the entire kernel shares a bitmap!
 - **Looking for specialized gadgets manually is impractical** even if we know what we are looking for
 - Needle in the haystack problem
- **kCFG** and **gadget search**
 - From our primitive **we know:**
 - Argument count
 - What is controlled within the function arguments
 - Desired execution primitive (decrement, increment,..)
 - **Lazarus** used an **undisclosed kCFG gadget** in **CVE-2024-21338**

kCFG constrained search



Parameter count → HLIL complexity threshold

Call filtering

Semantic
Code
Analysis

```
PS C:\Users\b33f\tools\binja-scripts\cfg-gadgets> python .\cfg.py C:\Windows\System32\drivers\acpi.sys --number 1 --line_threshold 15
--calls_with_args ExFreePoolWithTag
[+] CFG Metadata
|_ _load_config_used: 0x1c00600a0
|_ guardFlags offset: 0x90
|_ guardFlags: 0x1051c500
|_ Block Size: 0x5

[>] GFIDS section found at address 0x1c00ad000 with size 3032 bytes.
|_ Found 606 entries

[+] Entry 523: 0x87540 Absolute Address: 0x1c0087540 Function: ACPIDereferencePnpLocationInterface (1 parameters)
|_ HLIL Line Count: 9
|_ Signature: void ACPIDereferencePnpLocationInterface(void* arg1)
|_ Flag: 00
|_ Function Body:
1c0087540 void ACPIDereferencePnpLocationInterface(void* arg1)
1c0087546: int32_t temp0 = *(arg1 + 0x30)
1c0087546: *(arg1 + 0x30) -= 1
1c008754d: if (temp0 == 1)
1c008754f: int64_t rax_1 = *(arg1 + 0x20)
1c0087556: if (rax_1 != 0)
1c008755c: rax_1(*(arg1 + 0x10))
1c0087567: ExFreePoolWithTag(arg1, 0)
1c0087578: return
```

kCFG rendered irrelevant

- Avast did not share the kCFG gadget used by Lazarus
 - Reason? Avoid "making exploitation too widely accessible"
 - <https://decoded.avast.io/janvojtesek/lazarus-and-the-fudmodule-rootkit-beyond-byovd-with-an-admin-to-kernel-zero-day/>
- It doesn't matter: Our method renders kCFG completely ineffective
 - Quickly search thousands of gadgets with specific requirements for any scenario by querying the semantic structure of valid kCFG targets
 - kCFG payloads can be generated automatically based on obtained primitive
 - trigger vulnerability multiple times to chain calls
 - or use "dispatcher" gadgets
 - Bypass In-The-Wild (ITW) detections
 - Vary exploit chain patterns
 - Switch up call stack



Case Study CVE-2024-30089 Pool Spray

- Vulnerability exploited by chompie at Pwn2Own 2024 in Vancouver
 - Logic bug leading to a use after free with a very tight window before object reuse
 - use UAF to overwrite vtable pointer and gain kernel execution
 - Named pipe unbuffered 'IoSB' allocations
 - ExAllocatePoolWithTag (Win 10)
 - ExAllocatePool2 (Win11 - deprecated)
- Vulnerability complexity
 - ExAllocatePool2 zero initializes by default
 - Losing race in critical "zero initialized" state will cause NULL pointer dereference and result in BSOD
- Automation to the rescue
 - Two options to avoid zero initialization
 - Allocations with POOL_FLAG_UNINITIALIZED
 - Allocations using deprecated API
 - Object needs to be initialized quickly after allocation to have a chance of winning race -- restrict results with direct copies to returned buffer within < 3 LoC (HLIL_VAR_INIT)
 - Search space includes all Kernel modules/drivers reachable from user space, medium integrity
- Search methodology found an object that met these exact requirements and was used to successfully spray in the tight condition window
 - W32PIDLOCK!win32kbase

Case Study CVE-2024-30089 kCFG

- Evolving requirements for gadgets
 - Not all gadgets are made equal!
 - Initial exploit used kCFG gadget within the module of the originating call
 - Second argument **not controllable**
 - **Reliance on volatile register being odd** (reg & 1 == 1) for double-free primitive
 - About **50%** effective before **BSOD**
- Search scope expanded to all valid kCFG targets
 - If the **object reference count** is **set to 1**, that guarantees that the object will be freed in vtable release \Rightarrow **double-free primitive**
 - Using our **targeted search** with some custom conditions resulted in the gadget being found in **about 30 minutes** (VM 8GB RAM)
 - Increased exploit reliability to **100%**

```
# Iterate over the HLIL instructions and increment the line count
try:
    for x in hlil_function.instructions:
        if str(x) == "*this + 0x18) = 1":
            return True
        if str(x) == "*arg1 + 0x18) = 1":
            return True

except Exception as e:
    return False
```



```
[+] Entry 228: 0x1a1f0 Absolute Address: 0x1c001a1f0 Function: CTokenBase::Pending (lines: 3)
|_ Signature: void ?Pending@CTokenBase@@UEAAXXZ(CTokenBase* this)
|_ HLIL Gadget Found: *(this + 0x18) = 1
```

Variant analysis

- We analysed a vulnerability
 - What happened here?
 - How did this issue become exploitable?
 - What erroneous thought processes did the developer likely follow?
- We want to investigate, has this happened before
 - Once we understand the vulnerability we can break it down into its HLIL logical component parts
 - Run a semantic code search on HLIL - like a CodeQL query but for closed source binaries
 - When some combination of features exists elsewhere, we can say that they reproduce the same type of issue



Case-study CVE-2023-21768

- Root Cause Analysis (RCA)

- Introduction of **conditional statement** referring to **PreviousMode value** (KernelMode or UserMode)
- **Both conditions** execute the **same code**, else condition **missing ProbeForWrite** for UserMode case
- **Hard to identify** ↗ small mistake, large function, nondescript variable names

- Distilling code

- Identify all **conditional code** blocks with comparisons for **1 or 0**
- **Both conditions** execute **identical** code
- Code block should **assign value** or **call using value argument**
- Place **length limit** on the conditional **HLIL**

```
1c006fb77  
1c006fba4  
1c006fb77  
1c006fb81  
1c006fb81
```

```
if (l.b == 0)  
    *(ioRequestData + 0x18) = entriesRemoved  
else  
    *(ioRequestData + 0x18) = entriesRemoved
```

[+] Identical blocks found in function AfdNotifyRemoveIoCompletion at address 0x1c006fb77
|_ Condition: Length.b == 0, Block Length: 0
|_ True Block:
 **(arg3 + 0x18) = var_304
|_ False Block:
 **(arg3 + 0x18) = var_304



What about **other Kernel modules**

Talking a closer look

```
[*] Analyzing C:\Windows\System32\drivers\ks.sys...  
[+] Identical blocks found in function SerializePropertySet at address 0x1c00346a8  
|_ Condition: Irp->RequestorMode == 0, Block Length: 1  
|_ True Block:  
|   memcpy(rax_2, rdx_2, r14_1)  
|_ False Block:  
|   memcpy(rax_2, rdx_2, r14_1)  
  
[+] Identical blocks found in function UnserializePropertySet at address 0x1c0034986  
|_ Condition: Irp->RequestorMode == 0, Block Length: 1  
|_ True Block:  
|   memcpy(rax_2, rdx_1, r15_1)  
|_ False Block:  
|   memcpy(rax_2, rdx_1, r15_1)
```



```
VOID pType3InputBuffer = *(uint64_t*)((char*)pStack->Parameters + 0x18);  
  
if (!pIrp->RequestorMode)  
|   memcpy(P, pType3InputBuffer, NumberOfBytes);  
else  
|   memcpy(P, pType3InputBuffer, NumberOfBytes);
```

- A system-wide variant scan returned results
 - We took a closer look at some of the findings, including this one
 - At the time we, *incorrectly*, evaluated this as non-exploitable 😕
 - Later this was assigned **CVE-2024-30084**
- This was exploited by Devcore at Pwn2Own 2024 in Vancouver
 - Angelboy (@scwuaptx) combined **CVE-2024-30084** and **CVE-2024-35250** to achieve LPE
 - <https://devco.re/blog/2024/08/23/streaming-vulnerabilities-from-windows-kernel-proxying-to-kernel-part1-en/>

TOCTOU ⇢ ks.sys



Maintenance is ~~not~~ sexy

- Maintenance is **not super exciting**
 - The exploit **works on my machine**
 - The exploit **doesn't crash my target**
 - The exploit **targets my version**
- **Compatibility** and **stability** are essential
 - We really **want to do cool hacking shit** but it is important to remember that **the product has to be useful**
 - Maintenance covers many areas like VM automation, automated application deployment and conditional code updates
 - Common code update tasks include **symbol resolution** and **gadget portability**

Symbol and structure resolution

```
PS C:\Users\b33f\tools\binja-scripts\symbol-searcher> python .\search-symbol.py C:\Windows\System32\drivers\acpi
-p SOURCE_DESCRIPTOR
[!] No symbols matched the partial name.

[+] Matched Types:
+-----+-----+
| Type Name | Type Definition |
+-----+-----+
| CM_FULL_RESOURCE_DESCRIPTOR | struct _CM_FULL_RESOURCE_DESCRIPTOR |
+-----+-----+
| CM_PARTIAL_RESOURCE_DESCRIPTOR | struct _CM_PARTIAL_RESOURCE_DESCRIPTOR |
+-----+-----+
| IO_RESOURCE_DESCRIPTOR | struct _IO_RESOURCE_DESCRIPTOR |
+-----+-----+
| PCM_FULL_RESOURCE_DESCRIPTOR | struct _CM_FULL_RESOURCE_DESCRIPTOR* |
+-----+-----+
| PCM_PARTIAL_RESOURCE_DESCRIPTOR | struct _CM_PARTIAL_RESOURCE_DESCRIPTOR* |
+-----+-----+
| PHYSICAL_COUNTER_RESOURCE_DESCRIPTOR | struct _PHYSICAL_COUNTER_RESOURCE_DESCRIPTOR |
+-----+-----+
| PHYSICAL_COUNTER_RESOURCE_DESCRIPTOR_TYPE | enum _PHYSICAL_COUNTER_RESOURCE_DESCRIPTOR_TYPE |
+-----+-----+
| PIO_RESOURCE_DESCRIPTOR | struct _IO_RESOURCE_DESCRIPTOR* |
+-----+-----+
```



```
[+] Searching for unique symbol 'AcpiWrapperReadConfig'

Symbol Type: FunctionSymbol
Symbol Address: 0x1C0001810
Symbol Name: AcpiWrapperReadConfig
Function Name: AcpiWrapperReadConfig
Return Type: int64_t
Arguments:
  Arg 1: int32_t arg1
  Arg 2: int32_t arg2
```

[+] Searching for unique type '_WHEA_ERROR_SEVERITY':

Type Name: _WHEA_ERROR_SEVERITY

Type Definition: Enumeration (int32_t)

Enumerators:

Enumerator	Value
WheaErrSevRecoverable	0x0
WheaErrSevFatal	0x1
WheaErrSevCorrected	0x2
WheaErrSevInformational	0x3

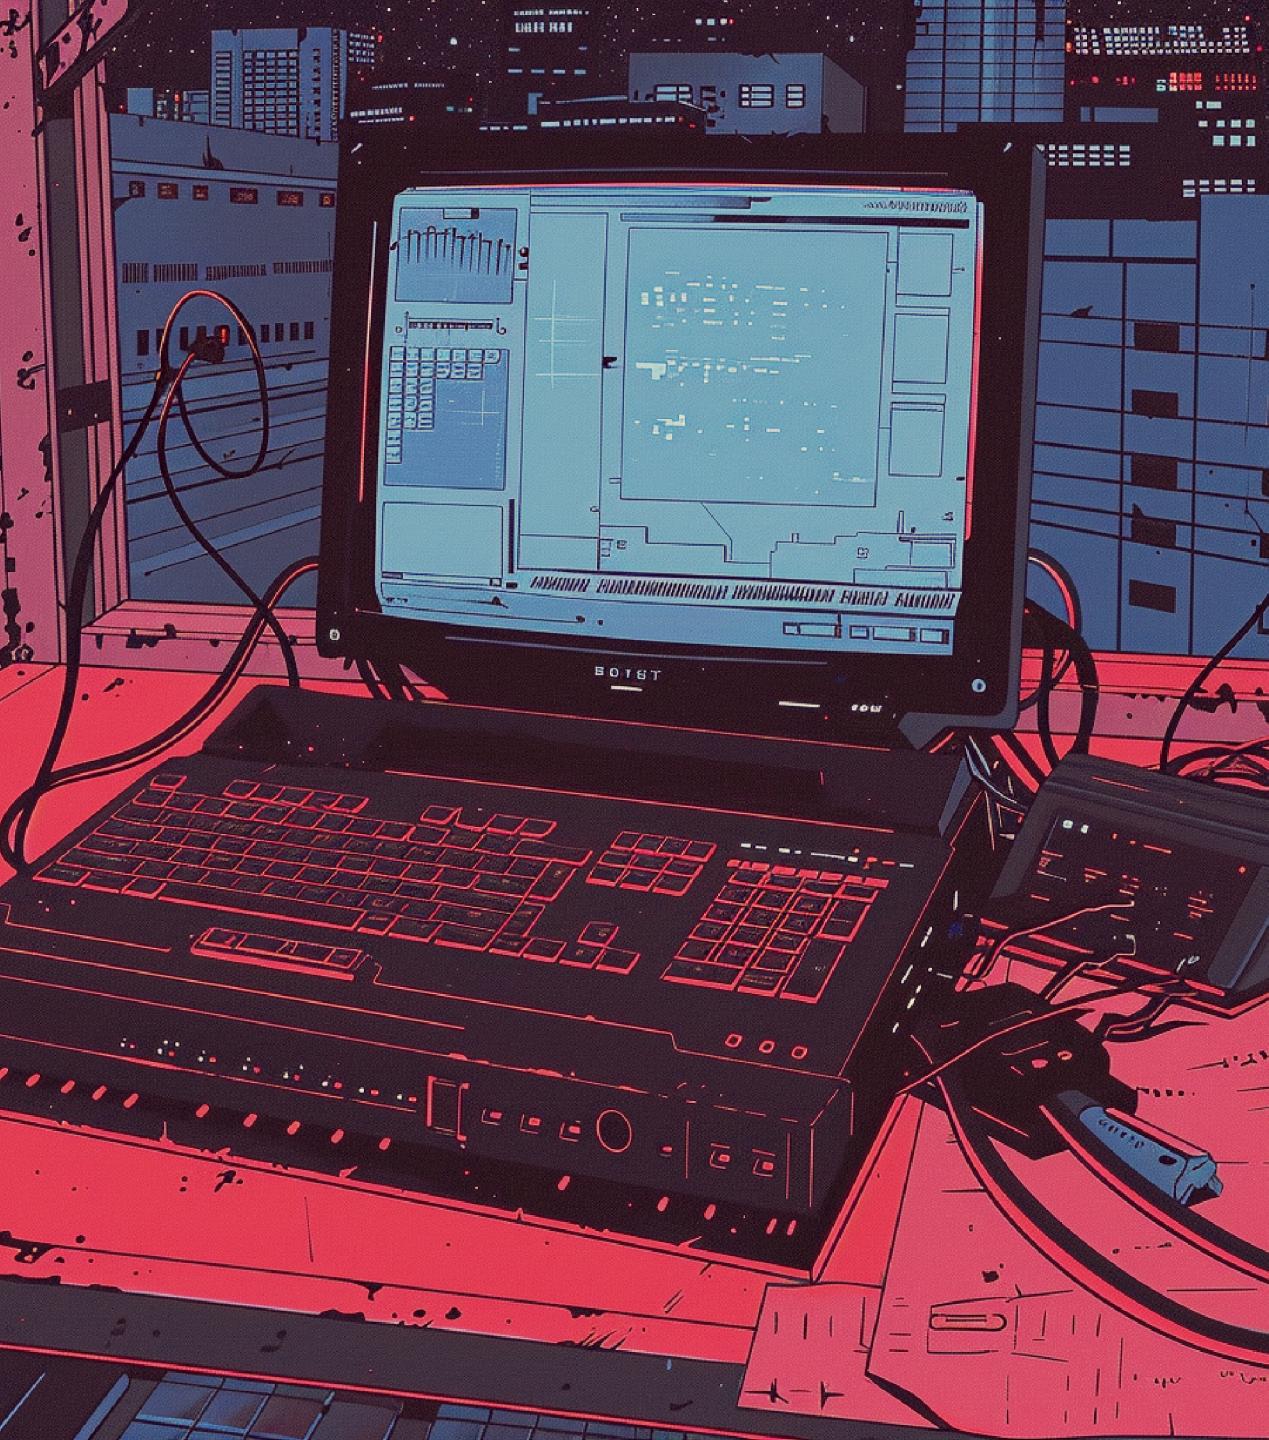
[+] Searching for unique type '_WHEA_ERROR_RECORD_HEADER':

Type Name: _WHEA_ERROR_RECORD_HEADER

Type Definition: Structure

Members:

Offset	Type	Name
0x0	ULONG	Signature
0x4	WHEA_REVISION	Revision
0x8	ULONG	SignatureEnd
0xC	USHORT	SectionCount
0x10	WHEA_ERROR_SEVERITY	Severity
0x14	WHEA_ERROR_RECORD_HEADER_VALIDBITS	ValidBits
0x18	ULONG	Length
0x20	WHEA_TIMESTAMP	Timestamp
0x28	GUID	PlatformId
0x38	GUID	PartitionId
0x48	GUID	CreatorId
0x58	GUID	NotifyType
0x68	ULONGLONG	RecordId
0x70	WHEA_ERROR_RECORD_HEADER_FLAGS	Flags
0x78	WHEA_PERSISTENCE_INFO	PersistenceInfo
0x80	UCHAR[0xc]	Reserved



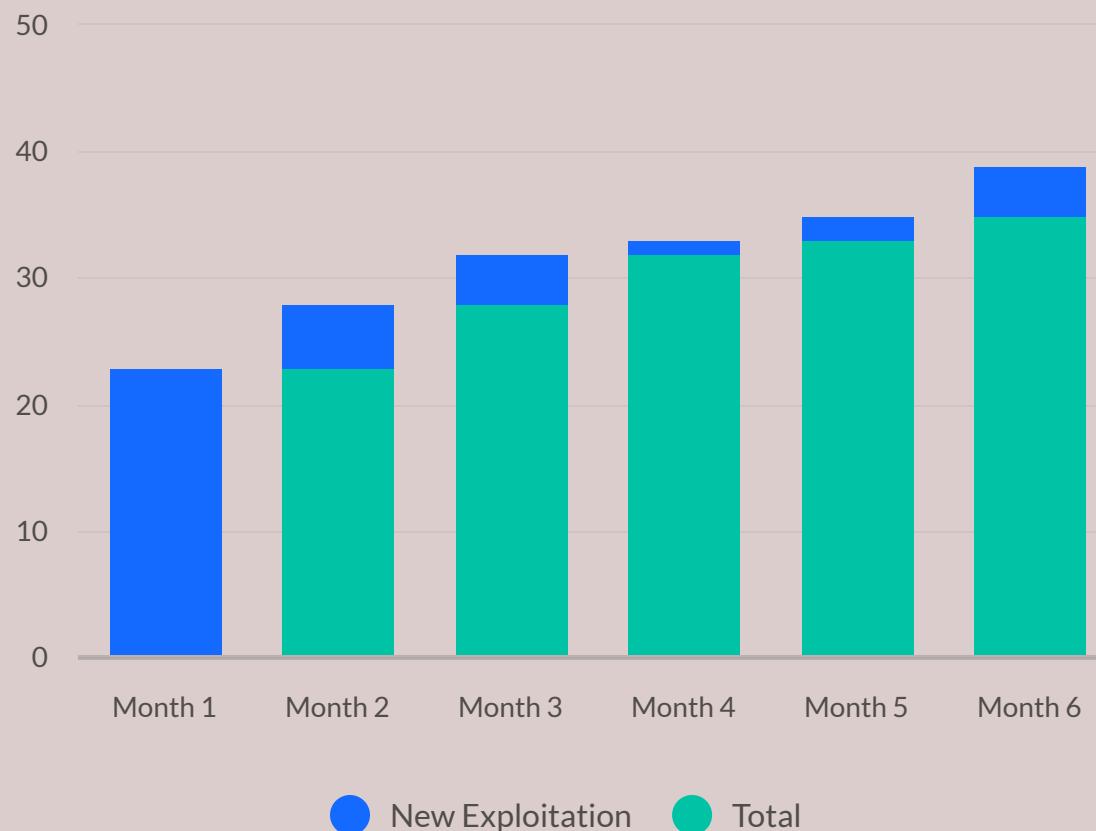
Conclusions and Future-Work

Life on the loop

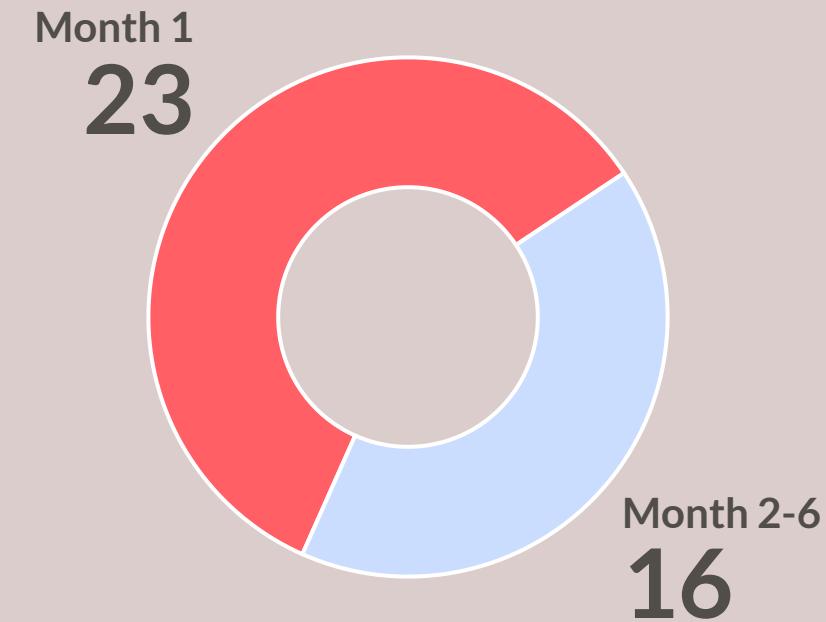
Time-to-Exploit Trends

Most exploitation happens in the first month after disclosure, [data from Mandiant](#)

Breakdown by Month



Total Distribution



If it ain't broke... still fix it!

- Many compelling reasons to build a library of swappable exploit objects & primitives
 - In-The-Wild (ITW) detections
 - Even commercial EDRs can detect mainstream exploitation methods
 - CrowdStrike knows about your pipe sprays 
- Staying ahead of the curve reduces future work
 - Future mitigations will continue to make exploitation harder
 - Beneficial to have multiple back up methods that do the same thing

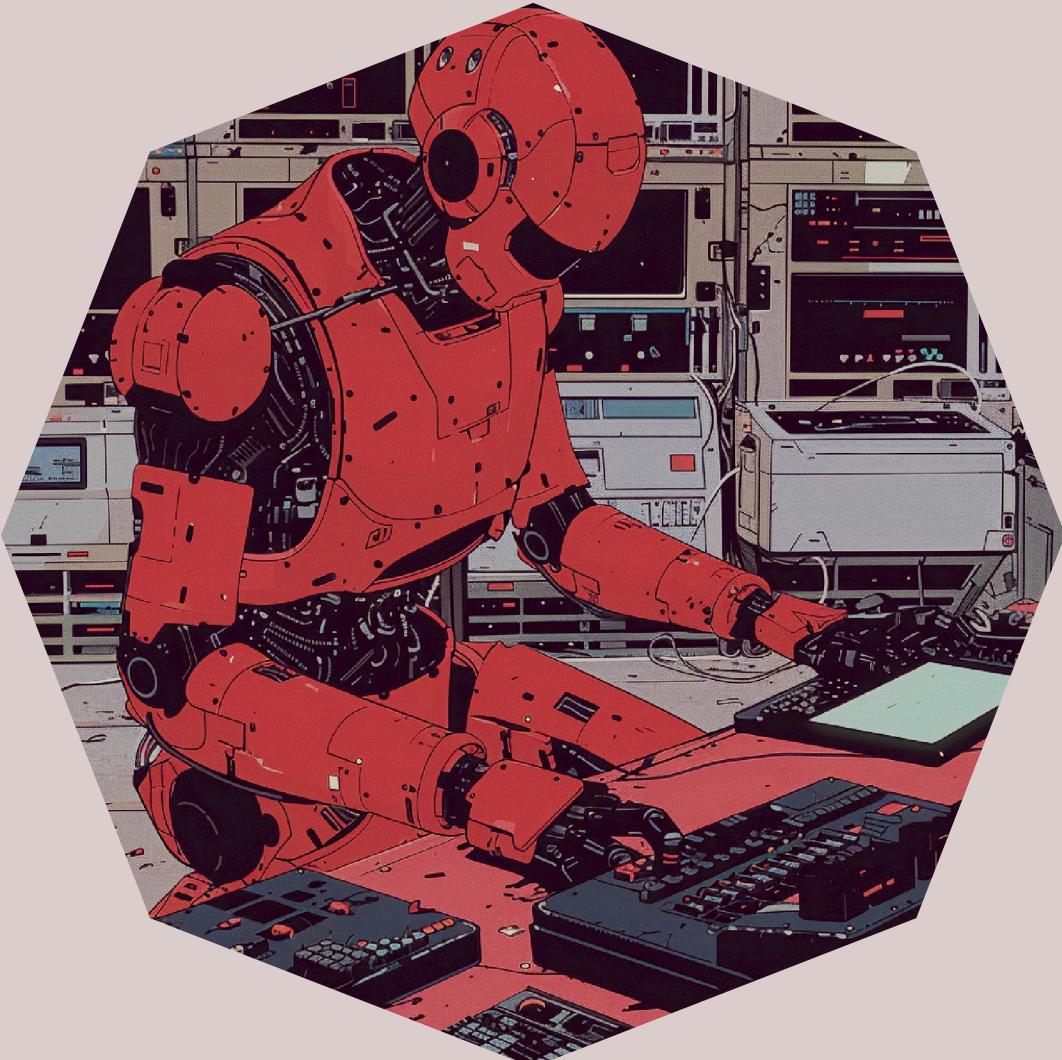


Conclusions

- Vulnerability research and exploit development is **hard work**
 - Even though suffering is a requirement, **you shouldn't suffer needlessly**
 - **Complex** and **deterministic** tasks should be **automated** to reduce manual labour
- **Training AI agents**
 - Training specialized agents is key, **layer specialized knowledge** on top of **frontier models**
 - Models are already very smart for general use cases but not yet in for this specialized use case
 - Model "**smartness**" is the **only bottleneck that matters**
 - **AI** will play a crucial role in **NatSec** and **VR/RE**
- **Program Analysis** should **power** automated VR and be the **backbone of AI reasoning**



Problem statements



- Train LLM's to reason better about Root Cause Analysis (RCA) and PoC generation
 - Multi-agent workflow
 - Planner, Reasoner, Questioner
- Future Workflow Improvements
 - Allocation search doesn't take into account object persistence
 - Can try to solve by looking for free calls in non failure paths
 - Finetuning, RAG and vector storage for special characteristics combined with extensive tool use (e.g. SMT solvers)
 - Models fail on complex scenarios and require hand-holding, partially engineering problem, partially fundamental "smartness" problem
 - Reaching vulnerable code doesn't mean the vulnerability is triggered, specific conditions must be met

Questions?

